

Homework 1

Luis Vitor Zerkowski - 14895730

September 19, 2023

1 Multivariate Calculus

1.a

Since $x \in \mathbb{R}^m$, and we are applying the sigmoid function to each one of the vector's entries, we have:

$$\sigma(x) = \begin{bmatrix} \sigma(x_1) \\ \vdots \\ \sigma(x_m) \end{bmatrix}$$

Thus, getting the gradient of the sigmoid function applied to the vector, $\nabla_x \sigma(x)$, is just a matter of applying the derivative of the sigmoid function to each one of the entries.

We then start by computing the derivative of the sigmoid function:

$$\frac{\partial \sigma(x)}{\partial x} = \frac{\partial}{\partial x} \frac{1}{1 + e^{-x}} =$$

By the quotient rule, the sum rule on $(1 + e^{-x})$ and also the chain rule on $(1 + e^{-x})$:

$$\begin{aligned} &= \frac{0 - (-1)e^{-x}}{(1 + e^{-x})^2} = \frac{e^{-x}}{(1 + e^{-x})^2} = \\ &= \frac{1}{1 + e^{-x}} \frac{e^{-x}}{1 + e^{-x}} = \sigma(x) \frac{e^{-x} + 1 - 1}{1 + e^{-x}} = \\ &= \sigma(x) \left(\frac{1 + e^{-x}}{1 + e^{-x}} - \frac{1}{1 + e^{-x}} \right) = \sigma(x)(1 - \sigma(x)) \end{aligned}$$

Having the derivative of the sigmoid function in hands, we know that:

$$\nabla_x \sigma(x) = \begin{bmatrix} \frac{\partial \sigma(x_1)}{x_1} \\ \vdots \\ \frac{\partial \sigma(x_m)}{x_m} \end{bmatrix} = \begin{bmatrix} \sigma(x_1)(1 - \sigma(x_1)) \\ \vdots \\ \sigma(x_m)(1 - \sigma(x_m)) \end{bmatrix}$$

1.b

Let's start by doing the multiplication with index notation to make it easier to compute the derivatives. Since $X \in \mathbb{R}^{n \times n}$ and $w \in \mathbb{R}^{n \times 1}$, multiplying both should give us $Xw \in \mathbb{R}^{n \times 1}$. So we do:

$$[Xw]_i = \sum_{p=1}^n X_{ip} w_p$$

Applying the derivatives w.r.t w on the i -th row, we get:

$$\frac{\partial}{\partial w_j} [Xw]_i = \frac{\partial}{\partial w_j} \sum_{p=1}^n X_{ip} w_p =$$

By the sum rule and also using the fact that X_{ip} is just a constant:

$$= \sum_{p=1}^n X_{ip} \frac{\partial w_p}{\partial w_j}$$

We now end up with a Kronecker Delta:

$$\frac{\partial w_p}{\partial w_j} = \delta_{jp} = \begin{cases} 1, & \text{if } j = p \\ 0, & \text{otherwise} \end{cases}$$

And then we can do:

$$\sum_{p=1}^n X_{ip} \frac{\partial w_p}{\partial w_j} = \sum_{p=1}^n X_{ip} \delta_{jp} = X_{ij}$$

Thus the derivative $\frac{\partial}{\partial w} f$ with $f = Xw$ give us exactly $\frac{\partial}{\partial w} f = X$.

1.c

We start again by doing the multiplication with index notation to make it easier to compute the derivatives. Since $X \in \mathbb{R}^{n \times n}$ and $w \in \mathbb{R}^{n \times 1}$, the multiplication $w^T X w$ should give us a real number. So we do:

$$w^T X w = \sum_{p=1}^n \sum_{q=1}^n w_p X_{pq} w_q$$

Applying the derivatives w.r.t w , we get:

$$\frac{\partial w^T X w}{\partial w_i} = \frac{\partial}{\partial w_i} \sum_{p=1}^n \sum_{q=1}^n w_p X_{pq} w_q =$$

By the sum rule and also using the fact that X_{pq} is just a constant:

$$= \sum_{p=1}^n \sum_{q=1}^n X_{pq} \frac{\partial}{\partial w_i} w_p w_q =$$

By the product rule:

$$= \sum_{p=1}^n \sum_{q=1}^n X_{pq} \left(\frac{\partial w_p}{\partial w_i} w_q + w_p \frac{\partial w_q}{\partial w_i} \right) =$$

Using the Kronecker Delta once again:

$$\begin{aligned} &= \sum_{p=1}^n \sum_{q=1}^n X_{pq} (\delta_{ip} w_q + w_p \delta_{iq}) = \sum_{p=1}^n \sum_{q=1}^n (X_{pq} \delta_{ip} w_q + X_{pq} w_p \delta_{iq}) = \\ &= \sum_{q=1}^n X_{iq} w_q + \sum_{p=1}^n X_{pi} w_p = Xw + w^T X = \\ &= w^T X^T + w^T X = w^T (X^T + X) \end{aligned}$$

Thus the derivative $\frac{\partial}{\partial w} f$ with $f = w^T X w$ give us exactly $\frac{\partial}{\partial w} f = w^T (X^T + X)$.

1.d

We start again by doing the multiplication with index notation to make it easier to compute the derivatives. Let's also call $(x - As) = v$ to simplify the problem for now. Since $\Sigma^{-1} \in \mathbb{R}^{m \times m}$ and $v \in \mathbb{R}^{m \times 1}$, the multiplication $v^T \Sigma^{-1} v$ should give us a real number. So we do:

$$v^T \Sigma^{-1} v = \sum_{p=1}^m \sum_{q=1}^m v_p \Sigma_{pq}^{-1} v_q$$

Applying the derivatives w.r.t Σ^{-1} , we get:

$$\frac{\partial v^T \Sigma^{-1} v}{\partial \Sigma_{ij}^{-1}} = \frac{\partial}{\partial \Sigma_{ij}^{-1}} \sum_{p=1}^m \sum_{q=1}^m v_p \Sigma_{pq}^{-1} v_q =$$

By the sum rule and also using the fact that v_p and v_q are just constants:

$$= \sum_{p=1}^m \sum_{q=1}^m v_p v_q \frac{\partial}{\partial \Sigma_{ij}^{-1}} \Sigma_{pq}^{-1} =$$

And since $\frac{\partial}{\partial \Sigma_{ij}^{-1}} \Sigma_{pq}^{-1} = 1$ if $i = p, j = q$ and $\frac{\partial}{\partial \Sigma_{ij}^{-1}} \Sigma_{pq}^{-1} = 0$ otherwise, it follows:

$$= \sum_{p=1}^m \sum_{q=1}^m v_p v_q \frac{\partial}{\partial \Sigma_{ij}^{-1}} \Sigma_{pq}^{-1} = v_i v_j = (x - As)(x - As)^T$$

Thus the derivative $\frac{\partial}{\partial \Sigma^{-1}} (x - As)^T \Sigma^{-1} (x - As)$ give us exactly $(x - As)(x - As)^T$.

1.e

We start by writing $\varsigma(x)$ with $x \in \mathbb{R}^n$ in its vector form:

$$\varsigma(x) = \begin{bmatrix} \frac{e^{x_1}}{\sum_{j=1}^n e^{x_j}} \\ \vdots \\ \frac{e^{x_n}}{\sum_{j=1}^n e^{x_j}} \end{bmatrix}$$

Thus, getting the derivative of ς w.r.t. the vector x is just a matter of applying the partial derivatives to each one of the rows.

$$\frac{\partial}{\partial x} \varsigma(x) = \begin{bmatrix} \frac{\partial}{\partial x_1} \frac{e^{x_1}}{\sum_{j=1}^n e^{x_j}} & \cdots & \frac{\partial}{\partial x_n} \frac{e^{x_1}}{\sum_{j=1}^n e^{x_j}} \\ \vdots & \ddots & \vdots \\ \frac{\partial}{\partial x_1} \frac{e^{x_n}}{\sum_{j=1}^n e^{x_j}} & \cdots & \frac{\partial}{\partial x_n} \frac{e^{x_n}}{\sum_{j=1}^n e^{x_j}} \end{bmatrix}$$

We then start by computing the derivative of a cell to fill out the matrix afterwards. For x_i we have:

$$\begin{aligned} \frac{\partial}{\partial x_i} \varsigma(x)_i &= \frac{\partial}{\partial x_i} \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}} = \\ &= \frac{\partial}{\partial x_i} \frac{e^{x_i}}{e^{x_1} + \dots + e^{x_n}} = \end{aligned}$$

Applying the quotient rule to the fraction and then the sum rule to $(e^{x_1} + \dots + e^{x_n})$, we get:

$$\begin{aligned} &= \frac{e^{x_i}(e^{x_1} + \dots + e^{x_n}) - e^{x_i}(0 + \dots + e^{x_i} + \dots + 0)}{(e^{x_1} + \dots + e^{x_n})^2} = \frac{e^{x_i}(e^{x_1} + \dots + e^{x_n}) - e^{2x_i}}{(e^{x_1} + \dots + e^{x_n})^2} = \\ &= \frac{e^{x_i}}{(e^{x_1} + \dots + e^{x_n})} - \frac{e^{x_i}}{(e^{x_1} + \dots + e^{x_n})} \frac{e^{x_i}}{(e^{x_1} + \dots + e^{x_n})} = \varsigma(x)_i - \varsigma(x)_i^2 \end{aligned}$$

Now for x_k with $k \neq i$, we have:

$$\begin{aligned}\frac{\partial}{\partial x_k} \varsigma(x)_i &= \frac{\partial}{\partial x_k} \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}} = \\ &= \frac{\partial}{\partial x_k} \frac{e^{x_i}}{e^{x_1} + \dots + e^{x_n}} =\end{aligned}$$

Applying the quotient rule to the fraction and then the sum rule to $(e^{x_1} + \dots + e^{x_n})$, we get:

$$\begin{aligned}&= \frac{0(e^{x_1} + \dots + e^{x_n}) - e^{x_i}(0 + \dots + e^{x_k} + \dots + 0)}{(e^{x_1} + \dots + e^{x_n})^2} = \frac{-e^{x_i+x_k}}{(e^{x_1} + \dots + e^{x_n})^2} = \\ &= -\frac{e^{x_i}}{(e^{x_1} + \dots + e^{x_n})} \frac{e^{x_k}}{(e^{x_1} + \dots + e^{x_n})} = -\varsigma(x)_i \varsigma(x)_k\end{aligned}$$

With all the matrix entries in hand, we can write the derivative down as $\frac{\partial}{\partial x} \varsigma(x) = \text{diag}(\varsigma(x)) - \varsigma(x)\varsigma(x)^T$.

2 Full Analysis of a Distribution: Poisson Distribution

2.a

We first start by showing the distribution is well defined. So for an arbitrary $X = k$, we have:

$$p(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

And now if we show that every part of the expression is non-negative, we know that their multiplications or quotients will be non-negative.

Since $\lambda > 0$ by definition, we know $\lambda^k = \prod_{i=1}^k \lambda > 0$ because we are simply multiplying a positive number k times.

Now for $e^{-\lambda}$, we have $e^{-\lambda} = \frac{1}{e^\lambda}$ which is non-negative, because e^λ is non-negative and if we take the limit of the fraction with $\lambda > 0$ going to infinity, we get $\lim_{\lambda \rightarrow \infty} \frac{1}{e^\lambda} = 0$.

For the last part, we know $k! = k(k-1)\dots 1 > 0$, since k is the number of occurrences of an event and so it has to be a non-negative number and even if $k = 0$, we know $0! = 1$.

Since all the parts of the expression are non-negative, we know the whole expression is non-negative.

Now for showing that the Poisson distribution is normalized, we sum over all the possibilities of the events. So we have:

$$\sum_{k=0}^{\infty} p(X = k) = \sum_{k=0}^{\infty} \frac{\lambda^k e^{-\lambda}}{k!} =$$

But since $e^{-\lambda}$ does not depend on k , we can take it out of the summation as a constant multiplying all the terms.

$$= e^{-\lambda} \sum_{k=0}^{\infty} \frac{\lambda^k}{k!} =$$

Now using the given Taylor expansion for the exponential function, we have:

$$= e^{-\lambda} e^\lambda = 1$$

And we just showed that $\sum_{k=0}^{\infty} p(X = k) = 1$, so the distribution is normalized indeed.

2.b

We start by computing the mean of the distribution. So by definition, we have:

$$E[X] = \sum_{k=0}^{\infty} kp(X=k) = \sum_{k=0}^{\infty} k \frac{\lambda^k e^{-\lambda}}{k!} =$$

Evaluating the summation on $k = 0$ and summing the rest, we have:

$$= 0 + \sum_{k=1}^{\infty} k \frac{\lambda^k e^{-\lambda}}{k!} = \sum_{k=1}^{\infty} \frac{\lambda^k e^{-\lambda}}{(k-1)!} =$$

Using again the fact that $e^{-\lambda}$ is a constant to the summation over k and splitting λ^k in a convenient way, we have:

$$= e^{-\lambda} \sum_{k=1}^{\infty} \frac{\lambda \lambda^{k-1}}{(k-1)!} = e^{-\lambda} \lambda \sum_{k=1}^{\infty} \frac{\lambda^{k-1}}{(k-1)!} =$$

Now changing variables $j = k - 1$ and using the given Taylor expansion of the exponential function again, we get:

$$= e^{-\lambda} \lambda \sum_{j=0}^{\infty} \frac{\lambda^j}{j!} = e^{-\lambda} \lambda e^{\lambda} = \lambda$$

So the mean of the Poisson distribution is $E[X] = \lambda$.

Now for the variance, we start again by the definition:

$$\sigma^2(X) = E[X^2] - E[X]^2$$

Since we have $E[x]$, we know that $E[x]^2 = \lambda^2$, so we just need to compute $E[X^2]$:

$$E[X^2] = \sum_{k=0}^{\infty} k^2 p(X=k) = \sum_{k=0}^{\infty} k^2 \frac{\lambda^k e^{-\lambda}}{k!} =$$

Separating $e^{-\lambda}$ from the summation again, splitting the first term, $k = 0$, of the summation and also splitting λ^k , we have:

$$= e^{-\lambda} (0 + \sum_{k=1}^{\infty} k \frac{\lambda \lambda^{k-1}}{(k-1)!}) = e^{-\lambda} \lambda (\sum_{k=1}^{\infty} k \frac{\lambda^{k-1}}{(k-1)!}) =$$

Manipulating the expression to get a $(k-1)$ term, we have:

$$= e^{-\lambda} \lambda \sum_{k=1}^{\infty} ((k-1) \frac{\lambda^{k-1}}{(k-1)!} + \frac{\lambda^{k-1}}{(k-1)!}) = e^{-\lambda} \lambda (\sum_{k=1}^{\infty} (k-1) \frac{\lambda^{k-1}}{(k-1)!} + \sum_{k=1}^{\infty} \frac{\lambda^{k-1}}{(k-1)!}) =$$

Now computing the next term, $k = 1$, of the first summation, $\sum_{k=1}^{\infty} (k-1) \frac{\lambda^{k-1}}{(k-1)!}$, we get:

$$= e^{-\lambda} \lambda (0 + \sum_{k=2}^{\infty} (k-1) \frac{\lambda^{k-1}}{(k-1)!} + \sum_{k=1}^{\infty} \frac{\lambda^{k-1}}{(k-1)!}) = e^{-\lambda} \lambda (\lambda \sum_{k=2}^{\infty} \frac{\lambda^{k-2}}{(k-2)!} + \sum_{k=1}^{\infty} \frac{\lambda^{k-1}}{(k-1)!}) =$$

Changing variables $j = k - 2$ and $z = k - 1$, and finally using the given Taylor expansion of the exponential function again, we have:

$$= e^{-\lambda} \lambda (\lambda \sum_{j=0}^{\infty} \frac{\lambda^j}{j!} + \sum_{z=0}^{\infty} \frac{\lambda^z}{z!}) = e^{-\lambda} \lambda (\lambda e^{\lambda} + e^{\lambda}) = \lambda^2 + \lambda$$

Going back to the definition of variance, we get:

$$\sigma^2(X) = E[X^2] - E[X]^2 = \lambda^2 + \lambda - \lambda^2 = \lambda$$

So the variance of the Poisson distribution is $\sigma^2(X) = \lambda$.

2.c

Given a fixed time interval Δt , the Poisson distribution changes a bit. Normally λ refers to the average rate of occurrences of an event per unit time, but now we are working with a specific time interval, which leads to:

$$p(X = k) = \frac{(\lambda\Delta t)^k e^{-\lambda\Delta t}}{k!}$$

For the likelihood of the events, we have to compute $p(k_1, \dots, k_N | \lambda)$, but since they are i.i.d, we have:

$$p(k_1, \dots, k_N | \lambda) = \prod_{i=1}^N p(k_i) = \prod_{i=1}^N \frac{(\lambda\Delta t)^{k_i} e^{-\lambda\Delta t}}{k_i!}$$

Writing the log-likelihood, thus, gives us:

$$\begin{aligned} \log \prod_{i=1}^N \frac{(\lambda\Delta t)^{k_i} e^{-\lambda\Delta t}}{k_i!} &= \sum_{i=1}^N \log \left(\frac{(\lambda\Delta t)^{k_i} e^{-\lambda\Delta t}}{k_i!} \right) = \\ &= \sum_{i=1}^N \log((\lambda\Delta t)^{k_i} e^{-\lambda\Delta t}) - \sum_{i=1}^N \log(k_i!) = \sum_{i=1}^N \log(\lambda\Delta t)^{k_i} + \sum_{i=1}^N -\lambda\Delta t - \sum_{i=1}^N \log(k_i!) = \\ &= \sum_{i=1}^N k_i \log(\lambda\Delta t) - \sum_{i=1}^N \log(k_i!) - N\lambda\Delta t \end{aligned}$$

So the log-likelihood of the events assuming i.i.d and that all events are Poisson processes measured in a fixed interval time Δt is given by $\sum_{i=1}^N k_i \log(\lambda\Delta t) - \sum_{i=1}^N \log(k_i!) - N\lambda\Delta t$.

2.d

For this exercise we start by taking log-likelihood derivative w.r.t λ and equal to zero:

$$\lambda_{ML} = \arg \max_{\lambda} \sum_{i=1}^N k_i \log(\lambda\Delta t) - \sum_{i=1}^N \log(k_i!) - N\lambda\Delta t$$

Which leads to:

$$\frac{\partial}{\partial \lambda} \left(\sum_{i=1}^N k_i \log(\lambda\Delta t) - \sum_{i=1}^N \log(k_i!) - N\lambda\Delta t \right) = 0$$

By the sum rule, we can differentiate each part of the equation separately and then put them all together. So we have:

$$\begin{aligned} \frac{\partial}{\partial \lambda} \sum_{i=1}^N k_i \log(\lambda\Delta t) &= \sum_{i=1}^N k_i \frac{1}{\lambda\Delta t} \Delta t = \sum_{i=1}^N \frac{k_i}{\lambda} \\ \frac{\partial}{\partial \lambda} \sum_{i=1}^N \log(k_i!) &= 0 \\ \frac{\partial}{\partial \lambda} N\lambda\Delta t &= N\Delta t \end{aligned}$$

With all parts in hand, we continue:

$$\begin{aligned} \sum_{i=1}^N \frac{k_i}{\lambda} - N\Delta t &= 0 \implies \frac{1}{\lambda} \sum_{i=1}^N k_i = N\Delta t \implies \\ \implies \sum_{i=1}^N k_i &= N\Delta t \lambda \implies \lambda_{ML} = \frac{\sum_{i=1}^N k_i}{N\Delta t} \end{aligned}$$

So the maximum likelihood estimator is $\lambda_{ML} = \frac{\sum_{i=1}^N k_i}{N\Delta t}$.

2.e

Since λ_{ML} represents an estimate of the average amount of people entering a store every minute, we just need to multiply it by a factor of 60 to get an estimator for an hour. So we have $\lambda_h = 60 * \lambda_{ML}$.

2.f

To find λ_{MAP} , we start by detailing its expression:

$$\lambda_{MAP} = \arg \max_{\lambda} p(\lambda|k) =$$

From Bayes rule, we get:

$$= \arg \max_{\lambda} \frac{p(k|\lambda)p(\lambda)}{p(k)} = \arg \max_{\lambda} \frac{p(k|\lambda)p(\lambda)}{\int p(k|\lambda)p(\lambda)\partial\lambda}$$

As the question states, it can be hard to compute the integral in the denominator of the above expression. But because we are working with maximum a posteriori estimation, we are using the arg max operator on λ . This means we are looking through all the possible values of λ , and all of them will have the same common denominator, so we can simply ignore the integral and work on finding the λ that maximizes $p(k|\lambda)p(\lambda)$ without any effect on the estimator found. We then continue our computations:

$$\lambda_{MAP} = \arg \max_{\lambda} p(k|\lambda)p(\lambda) =$$

Using the log trick, we get:

$$= \arg \max_{\lambda} \log(p(k|\lambda)p(\lambda)) = \arg \max_{\lambda} \log p(k|\lambda) + \log p(\lambda)$$

Since we already have the log-likelihood, we just need to compute the expression for the prior $\log p(\lambda)$ before continuing with our λ_{MAP} computations. Considering the given prior distribution, we have:

$$\log p(\lambda) = \log p(\lambda|\alpha_1, \alpha_2) = \log \frac{\alpha_2^{\alpha_1}}{\Gamma(\alpha_1)} \lambda^{\alpha_1-1} e^{-\alpha_2\lambda} =$$

$$= \alpha_1 \log \alpha_2 - \log \Gamma(\alpha_1) + (\alpha_1 - 1) \log \lambda - \alpha_2 \lambda =$$

Expanding the Gamma function $\Gamma(x) = (x-1)!$:

$$= \alpha_1 \log \alpha_2 - \left(\sum_{i=1}^{\alpha_1-1} \log(\alpha_1 - i) \right) + (\alpha_1 - 1) \log \lambda - \alpha_2 \lambda$$

Going back to the original equation $\arg \max_{\lambda} \log p(k|\lambda) + \log p(\lambda)$, we get:

$$\begin{aligned} \arg \max_{\lambda} \log p(k|\lambda) + \log p(\lambda) &= \arg \max_{\lambda} \sum_{i=1}^N k_i \log \lambda - \sum_{i=1}^N \log(k_i!) - N\lambda + \alpha_1 \log \alpha_2 - \left(\sum_{i=1}^{\alpha_1-1} \log(\alpha_1 - i) \right) + (\alpha_1 - 1) \log \lambda - \alpha_2 \lambda = \\ &= \arg \max_{\lambda} \log \lambda \sum_{i=1}^N k_i - \sum_{i=1}^N \log(k_i!) + \alpha_1 \log \alpha_2 - \left(\sum_{i=1}^{\alpha_1} \log(\alpha_1 - i) \right) + (\alpha_1 - 1) \log \lambda - (N + \alpha_2)\lambda = \\ &= \arg \max_{\lambda} \left(\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1 \right) \log \lambda - (N + \alpha_2)\lambda - \sum_{i=1}^N \log(k_i!) + \alpha_1 \log \alpha_2 - \left(\sum_{i=1}^{\alpha_1-1} \log(\alpha_1 - i) \right) \end{aligned}$$

But to get to the value that maximizes the above expression w.r.t λ , we would take its derivative and equal it to zero, which would zero out all the terms that do not depend on λ . Therefore, we can rewrite the expression with only the terms that depend on λ without compromising the maximum a posteriori estimation:

$$\lambda_{MAP} = \arg \max_{\lambda} \left(\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1 \right) \log \lambda - (N + \alpha_2)\lambda$$

Exactly what was expected.

2.g

To find the MAP estimator we need to take the derivative of the expression we got on the last question w.r.t λ and equal it to zero. So we have:

$$\frac{\partial}{\partial \lambda} \left(\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1 \right) \log \lambda - (N + \alpha_2) \lambda = 0$$

By the sum rule:

$$\frac{\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1}{\lambda} - (N + \alpha_2) = 0$$

$$\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1 = (N + \alpha_2) \lambda$$

$$\lambda_{MAP} = \frac{\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1}{(N + \alpha_2)}$$

Thus, the MAP estimator is $\lambda_{MAP} = \frac{\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1}{(N + \alpha_2)}$.

2.h

Since we are only interested in showing the posterior distribution is also a Gamma distribution, we can ignore the evidence, which works as a normalizing constant, and compute $p(\lambda|k) \propto p(k|\lambda)p(\lambda)$:

$$\begin{aligned} p(k|\lambda)p(\lambda) &= \prod_{i=1}^N \frac{\lambda^{k_i} e^{-\lambda}}{k_i!} \frac{\alpha_2^{\alpha_1}}{\Gamma(\alpha_1)} \lambda^{\alpha_1-1} e^{-\alpha_2 \lambda} = \\ &= \frac{\lambda^{\sum_{i=1}^N k_i} e^{-N\lambda}}{\prod_{i=1}^N k_i!} \frac{\alpha_2^{\alpha_1}}{\Gamma(\alpha_1)} \lambda^{\alpha_1-1} e^{-\alpha_2 \lambda} \propto \lambda^{\left(\sum_{i=1}^N k_i \right) + \alpha_1 - 1} e^{-(N + \alpha_2) \lambda} \end{aligned}$$

And from the last expression, we observe that $p(k|\lambda)p(\lambda)$ is indeed proportional to a Gamma distribution with parameters $\alpha'_1 = (\sum_{i=1}^N k_i) + \alpha_1$ and $\alpha'_2 = N + \alpha_2$, so $p(\lambda|k)$ is indeed proportional to a Gamma distribution.

3 General Multiple Outputs Linear Regression

3.a

We can inspect the dimensions of the parameter W by understanding the dimensions of the other parameters around it. We know that $y(x, W) \in \mathbb{R}^{K \times 1}$, since we want to predict an output for each one of the K targets. We also know that $\phi(x)$ is an M -dimensional vector, so $\phi(x) \in \mathbb{R}^{M \times 1}$. By using these facts and the given equation $y(x, W) = W^T \phi(x)$, we conclude that $W^T \in \mathbb{R}^{K \times M}$, which then leads to the answer $W \in \mathbb{R}^{M \times K}$.

3.b

For each target vector, we have:

$$p(t_i|W, \Sigma) = N(t_i|y(x_i, W), \Sigma) = \frac{1}{\sqrt{(2\pi)^K \det \Sigma}} e^{\frac{-1}{2}(t_i - y(x_i, W))^T \Sigma^{-1} (t_i - y(x_i, W))}$$

So the log-likelihood is given by:

$$\begin{aligned} \log p(t_i|W, \Sigma) &= \log \frac{1}{\sqrt{(2\pi)^K \det \Sigma}} e^{-\frac{1}{2}(t_i - y(x_i, W))^T \Sigma^{-1} (t_i - y(x_i, W))} = \\ &= \left(-\frac{1}{2}(t_i - y(x_i, W))^T \Sigma^{-1} (t_i - y(x_i, W)) \right) - \log \sqrt{(2\pi)^K \det \Sigma} = \end{aligned}$$

$$\begin{aligned}
& = \left(-\frac{1}{2}(t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)) \right) - \frac{1}{2} \log((2\pi)^K \det \Sigma) = \\
& = \left(-\frac{1}{2}(t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)) \right) - \frac{K}{2} \log(2\pi) - \frac{1}{2} \log \det \Sigma = \\
& = -\frac{1}{2}((t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)) + K \log 2\pi + \log \det \Sigma)
\end{aligned}$$

Now that we have the log-likelihood for each target, we can compute the log-likelihood for all N independent observations. It will turn out to be a sum of log terms, since we can multiply the likelihoods for all the independent observations. So we have:

$$\begin{aligned}
p(T|W, \Sigma) &= \prod_{i=1}^N N(t_i|y(x_i, W), \Sigma) \\
\log p(T|W, \Sigma) &= \sum_{i=1}^N \log N(t_i|y(x_i, W), \Sigma) =
\end{aligned}$$

And now using the previous results:

$$\begin{aligned}
& = \sum_{i=1}^N -\frac{1}{2}((t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)) + K \log 2\pi + \log \det \Sigma) = \\
& = -\frac{1}{2}(NK \log 2\pi + N \log \det \Sigma + \sum_{i=1}^N (t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)))
\end{aligned}$$

So the log-likelihood of T is given by $\log p(T|W, \Sigma) = -\frac{1}{2}(NK \log 2\pi + N \log \det \Sigma + \sum_{i=1}^N (t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)))$.

3.c

To find the maximum likelihood solution W_{ML} , we need to take the derivative w.r.t W of the log-likelihood function computed above and equal it to zero. We start by computing the derivative w.r.t W^T and proceed to prove that its equal to the transpose of the derivative w.r.t W . So we have:

$$\frac{\partial}{\partial W^T} \left(-\frac{1}{2}(NK \log 2\pi + N \log \det \Sigma + \sum_{i=1}^N (t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W))) \right) =$$

We remove the first two terms $\frac{NK}{2} \log 2\pi$ and $\frac{N}{2} \log \det \Sigma$ since their derivative w.r.t W^T is zero. So we get:

$$\frac{\partial}{\partial W^T} \left(-\frac{1}{2} \left(\sum_{i=1}^N (t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)) \right) \right) =$$

And now by the sum rule:

$$-\frac{1}{2} \left(\sum_{i=1}^N \frac{\partial}{\partial W^T} (t_i - y(x_i, W))^T \Sigma^{-1}(t_i - y(x_i, W)) \right) = -\frac{1}{2} \left(\sum_{i=1}^N \frac{\partial}{\partial W^T} (t_i - W^T \phi(x_i))^T \Sigma^{-1}(t_i - W^T \phi(x_i)) \right) =$$

Making use of the hint given to us $\frac{\partial}{\partial A} (x - As)^T W (x - As) = -2W(x - As)s^T$ and the fact that the derivative w.r.t the transpose is equal to the transpose of the derivative w.r.t the original matrix , we have:

$$-\frac{1}{2} \left(\sum_{i=1}^N -2(\Sigma^{-1}(t_i - W^T \phi(x_i)) \phi(x_i)^T)^T \right) = \sum_{i=1}^N \phi(x_i)(t_i^T - \phi(x_i)^T W) \Sigma^{-1} =$$

Transforming it to the matrix form, we get:

$$= \Phi^T (T - \Phi W) \Sigma^{-1} = \Phi^T T \Sigma^{-1} - \Phi^T \Phi W \Sigma^{-1}$$

Now equaling it to zero, we have:

$$\begin{aligned}\Phi^T T \Sigma^{-1} - \Phi^T \Phi W \Sigma^{-1} &= 0 \\ \Phi^T \Phi W \Sigma^{-1} &= \Phi^T T \Sigma^{-1}\end{aligned}$$

But we know that the covariance matrix is symmetric and positive definite, thus it has an inverse - that we've been using - and we can do:

$$\begin{aligned}\Phi^T \Phi W \Sigma^{-1} \Sigma &= \Phi^T T \Sigma^{-1} \Sigma \\ \Phi^T \Phi W &= \Phi^T T\end{aligned}$$

We also know that $\Phi^T \Phi$ should be invertible as long as we get linear independent x input vectors.

Assuming we do, we then have:

$$W_{ML} = (\Phi^T \Phi)^{-1} \Phi^T T$$

So now we have shown that the maximum likelihood solution is $W_{ML} = (\Phi^T \Phi)^{-1} \Phi^T T$ and thus is independent of the covariance matrix Σ . We just need to show that this result is valid by proving the derivative of the likelihood w.r.t W^T is equal to the transpose of the derivative of the likelihood w.r.t W :

$$-\frac{1}{2} \left(\sum_{i=1}^N \frac{\partial}{\partial W} (t_i - W^T \phi(x_i))^T \Sigma^{-1} (t_i - W^T \phi(x_i)) \right)$$

Focusing on the derivative term and expanding to then use index notation, we have:

$$\begin{aligned}\frac{\partial}{\partial W} (t_i - W^T \phi(x_i))^T \Sigma^{-1} (t_i - W^T \phi(x_i)) &= \frac{\partial}{\partial W} (t_i^T - \phi(x_i)^T W) \Sigma^{-1} (t_i - W^T \phi(x_i)) = \\ &= \frac{\partial}{\partial W} (t_i^T \Sigma^{-1} - \phi(x_i)^T W \Sigma^{-1}) (t_i - W^T \phi(x_i)) = \\ &= \frac{\partial}{\partial W} (t_i^T \Sigma^{-1} t_i - \phi(x_i)^T W \Sigma^{-1} t_i - t_i^T \Sigma^{-1} W^T \phi(x_i) + \phi(x_i)^T W \Sigma^{-1} W^T \phi(x_i))\end{aligned}$$

By the sum rule, we can compute each derivative separately:

1.

$$\frac{\partial}{\partial W} t_i^T \Sigma^{-1} t_i = 0$$

2.

$$\begin{aligned}\frac{\partial}{\partial W} \phi(x_i)^T W \Sigma^{-1} t_i &\implies \sum_{p=1}^M \sum_{q=1}^K \sum_{r=1}^K \frac{\partial}{\partial W_{m,n}} \phi(x_i)_p W_{p,q} \Sigma_{q,r}^{-1} (t_i)_r = \sum_{r=1}^K \phi(x_i)_m \Sigma_{n,r}^{-1} (t_i)_r \implies \\ &\implies \frac{\partial}{\partial W} \phi(x_i)^T W \Sigma^{-1} t_i = \phi(x_i) (\Sigma^{-1} t_i)^T\end{aligned}$$

3.

$$\begin{aligned}\frac{\partial}{\partial W} t_i^T \Sigma^{-1} W^T \phi(x_i) &\implies \sum_{p=1}^K \sum_{q=1}^M \sum_{r=1}^M \frac{\partial}{\partial W_{m,n}} (t_i)_p \Sigma_{p,q}^{-1} W_{q,r} \phi(x_i)_r = \sum_{p=1}^K (t_i)_p \Sigma_{p,m}^{-1} \phi(x_i)_n \implies \\ &\implies \frac{\partial}{\partial W} t_i^T \Sigma^{-1} W^T \phi(x_i) = (t_i^T \Sigma^{-1})^T \phi(x_i)^T = \phi(x_i) (\Sigma^{-1} t_i)^T\end{aligned}$$

4.

$$\frac{\partial}{\partial W} \phi(x_i)^T W \Sigma^{-1} W^T \phi(x_i) \implies \sum_{p=i}^M \sum_{q=1}^K \sum_{r=1}^K \sum_{s=1}^M \frac{\partial}{\partial W_{m,n}} \phi(x_i)_p W_{p,q} \Sigma_{q,r}^{-1} W_{r,s} \phi(x_i)_s =$$

By the product rule:

$$\begin{aligned}
&= \sum_{p=i}^M \sum_{q=1}^K \sum_{r=1}^K \sum_{s=1}^M (\phi(x_i)_m \Sigma_{n,r}^{-1} W_{r,s} \phi(x_i)_s + \phi(x_i)_p W_{p,q} \Sigma_{q,m}^{-1} \phi(x_s)_n) = \\
&= \sum_{r=1}^K \sum_{s=1}^M \phi(x_i)_m \Sigma_{n,r}^{-1} W_{r,s} \phi(x_i)_s + \sum_{p=i}^M \sum_{q=1}^K \phi(x_i)_p W_{p,q} \Sigma_{q,m}^{-1} \phi(x_s)_n \implies \\
\implies \frac{\partial}{\partial W} \phi(x_i)^T W \Sigma^{-1} W^T \phi(x_i) &= \phi(x_i) (\Sigma^{-1} W^T \phi(x_i))^T + \phi(x_i) \phi(x_i)^T W \Sigma^{-1} =
\end{aligned}$$

Using the symmetry of Σ and thus Σ^{-1} :

$$= \phi(x_i) \phi(x_i)^T W \Sigma^{-1} + \phi(x_i) \phi(x_i)^T W \Sigma^{-1} = 2\phi(x_i) \phi(x_i)^T W \Sigma^{-1}$$

Putting all the terms together, we have:

$$\begin{aligned}
\frac{\partial}{\partial W} (t_i - W^T \phi(x_i))^T \Sigma^{-1} (t_i - W^T \phi(x_i)) &= -\phi(x_i) (\Sigma^{-1} t_i)^T - \phi(x_i) (\Sigma^{-1} t_i)^T + 2\phi(x_i) \phi(x_i)^T W \Sigma^{-1} = \\
&= -2\phi(x_i) (\Sigma^{-1} t_i)^T + 2\phi(x_i) \phi(x_i)^T W \Sigma^{-1} = -2\phi(x_i) t_i^T \Sigma^{-1} + 2\phi(x_i) \phi(x_i)^T W \Sigma^{-1} = \\
&= -2\phi(x_i) (t_i^T - \phi(x_i)^T W) \Sigma^{-1}
\end{aligned}$$

And if we put the obtained equation back to the original expression, we get:

$$-\frac{1}{2} \left(\sum_{i=1}^N \frac{\partial}{\partial W} (t_i - W^T \phi(x_i))^T \Sigma^{-1} (t_i - W^T \phi(x_i)) \right) = \sum_{i=1}^N \phi(x_i) (t_i^T - \phi(x_i)^T W) \Sigma^{-1}$$

Exactly the expression we got by using the hint and thus proving that indeed the derivative of the likelihood w.r.t W^T is equal to the transpose of the derivative of the likelihood w.r.t W .

So we can finally state that $W_{ML} = (\Phi^T \Phi)^{-1} \Phi^T T$.

3.d

We repeat the same process but now taking the derivative w.r.t $\Omega = \Sigma^{-1}$:

$$\frac{\partial}{\partial \Omega} \left(-\frac{1}{2} (NK \log 2\pi + N \log \det \Omega^{-1} + \sum_{i=1}^N (t_i - y(x_i, W))^T \Omega (t_i - y(x_i, W))) \right) =$$

We cancel out $-\frac{NK}{2} \log 2\pi$ since it does not depend on Ω :

$$= \frac{\partial}{\partial \Omega} \left(-\frac{1}{2} (N \log \det \Omega^{-1} + \sum_{i=1}^N (t_i - y(x_i, W))^T \Omega (t_i - y(x_i, W))) \right)$$

Separating both parts and computing them, we have:

$$\frac{\partial}{\partial \Omega} \log \det \Omega^{-1} =$$

By the relationship $\det X^{-1} = \frac{1}{\det X}$, we get:

$$= \frac{\partial}{\partial \Omega} \log \frac{1}{\det \Omega} = \frac{\partial}{\partial \Omega} (-\log \det \Omega)$$

Now using the given identity, we have:

$$\frac{\partial}{\partial \Omega} \log \det \Omega^{-1} = -(\Omega^{-1})^T$$

And for the second part:

$$\frac{\partial}{\partial \Omega} \sum_{i=1}^N (t_i - y(x_i, W))^T \Omega (t_i - y(x_i, W)) = \frac{\partial}{\partial \Omega} \sum_{i=1}^N (t_i - W^T \phi(x_i))^T \Omega (t_i - W^T \phi(x_i)) =$$

Now to use the second identity that was given to us, we need to transform Ω into Ω^T . We can do that simply by taking the transpose, since we know that covariance matrix is symmetric and the inverse of a symmetric matrix is also symmetric. Therefore, changing the equation to get to use the identity and also by the sum rule, we get:

$$= \sum_{i=1}^N \frac{\partial}{\partial \Omega} (t_i - W^T \phi(x_i))^T \Omega^T (t_i - W^T \phi(x_i)) = \sum_{i=1}^N (t_i - W^T \phi(x_i)) (t_i - W^T \phi(x_i))^T$$

Now putting everything together, we have:

$$\begin{aligned} \frac{\partial}{\partial \Omega} \left(-\frac{1}{2} (N \log \det \Omega^{-1} + \sum_{i=1}^N (t_i - y(x_i, W))^T \Omega (t_i - y(x_i, W))) \right) &= \\ = -\frac{1}{2} (-N(\Omega^{-1})^T + \sum_{i=1}^N (t_i - W^T \phi(x_i)) (t_i - W^T \phi(x_i))^T) \end{aligned}$$

Finally equaling it zero:

$$-N(\Omega^{-1})^T + \sum_{i=1}^N (t_i - W^T \phi(x_i)) (t_i - W^T \phi(x_i))^T = 0$$

Going back to Σ , we have:

$$\begin{aligned} N\Sigma^T &= \sum_{i=1}^N (t_i - W^T \phi(x_i)) (t_i - W^T \phi(x_i))^T \\ \Sigma^T &= \frac{1}{N} \sum_{i=1}^N (t_i - W^T \phi(x_i)) (t_i - W^T \phi(x_i))^T \end{aligned}$$

And using the symmetry of covariance matrix, we get:

$$\Sigma_{ML} = \frac{1}{N} \sum_{i=1}^N (t_i - W_{ML}^T \phi(x_i)) (t_i - W_{ML}^T \phi(x_i))^T$$

So it's shown that the maximum likelihood solution for Σ is given by $\Sigma_{ML} = \frac{1}{N} \sum_{i=1}^N (t_i - W_{ML}^T \phi(x_i)) (t_i - W_{ML}^T \phi(x_i))^T$.

4 Bayesian Linear Regression

4.a

To obtain the reported posterior, a Gaussian distribution $N(w|\mu_0, \Sigma_0)$ has been used as the prior. This can be understood exactly by answering the hint question. If no data is observed, the posterior should look like $p(w|t) = N(w|\mu_N, \Sigma_N)$ with $N = 0$, so $\mu_N = \Sigma_0 \Sigma_0^{-1} \mu_0 = \mu_0$ and $\Sigma_N^{-1} = \Sigma_0^{-1}$.

4.b

When the precision β approaches zero it means we lost confidence on our model, so adding new information does not update the model parameters. This can also be seen mathematically by $\lim_{\beta \rightarrow 0} \Sigma_{N+1}^{-1} = \lim_{\beta \rightarrow 0} \Sigma_N^{-1} + \beta \phi_{N+1} \phi_{N+1}^T = \Sigma_N^{-1}$, which also leads to $\lim_{\beta \rightarrow 0} \mu_{N+1} = \lim_{\beta \rightarrow 0} \Sigma_N (\Sigma_N^{-1} \mu_N + \beta \phi_{N+1} t_{N+1}) = \mu_N$, so $\Sigma_{N+1} = \Sigma_N$ and $\mu_{N+1} = \mu_N$.

4.c

Using the fact that the posterior becomes the prior for a new posterior after a new observation, we can write the updated posterior as:

$$p(w|\Phi_{N+1}, t_{N+1}, \Sigma_0, \mu_0, \beta) = \frac{p(t_{N+1}|w, \Phi_{N+1}, \Sigma_0, \mu_0, \beta)p(w|\Phi_N, t_N, \Sigma_0, \mu_0, \beta)}{\int p(t_{N+1}|w, \Phi_{N+1}, \Sigma_0, \mu_0, \beta)p(w|\Phi_N, t_N, \Sigma_0, \mu_0, \beta)dw}$$